

prediction-and-analysis-using-svm

July 27, 2024

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

# Load the dataset
df = pd.read_csv('diabetes.csv')

# View the dimensions of the data
print("Dimensions of the data:", df.shape)

# Gather summary statistics
summary_stats = df.describe()
print(summary_stats)
```

Dimensions of the data: (768, 9)

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

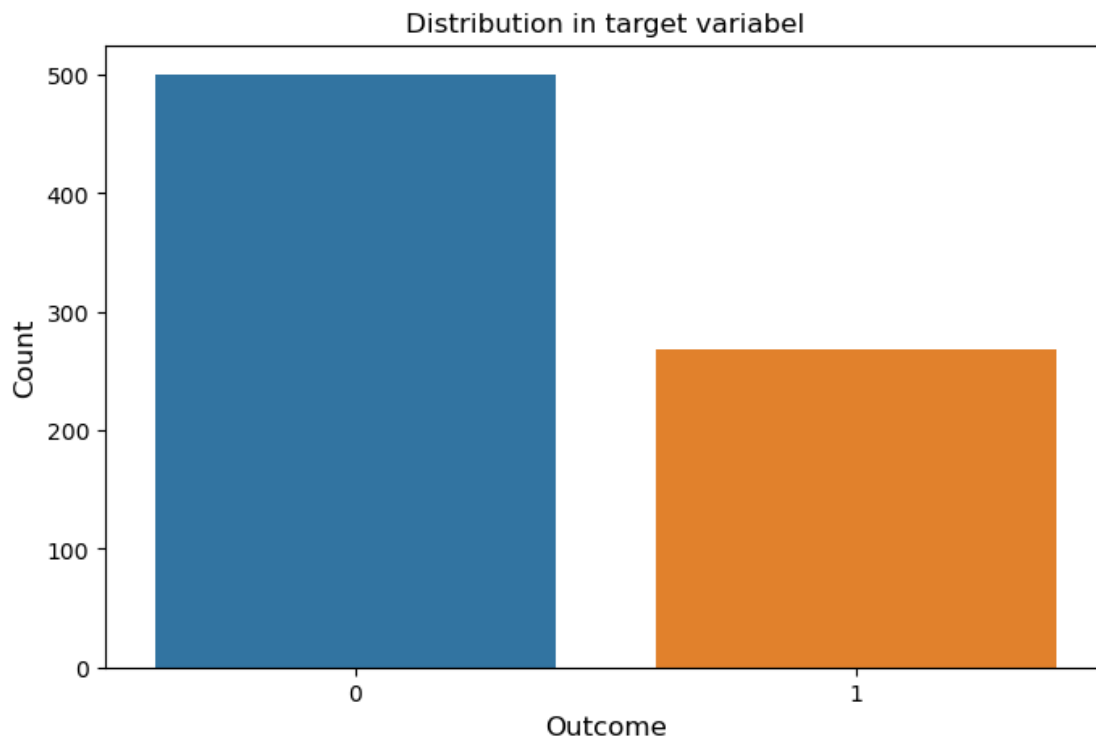
Observations:

1. There are missing values in columns , 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin'

2. The range of values varies over a large values. "Pregnancies" ranges from 0 to 17, "Glucose"
3. There is an imbalance in 'Outcome' column as the mean value shows only 34%. This shows that :
4. Huge difference in mean and median of 'Insulin' and 'SkinThickness' shows there is a skewness.

```
[2]: plt.figure(figsize = (8,5))
sns.countplot(x = df['Outcome'])
plt.xlabel('Outcome', size = 12)
plt.ylabel('Count', size = 12)
plt.title('Distribution in target variabel', size = 12)
```

```
[2]: Text(0.5, 1.0, 'Distribution in target variabel')
```



```
[3]: print(df.dtypes)
```

```
Pregnancies      int64
Glucose           int64
BloodPressure     int64
SkinThickness     int64
Insulin           int64
BMI               float64
DiabetesPedigreeFunction float64
```

```
Age                                int64
Outcome                           int64
dtype: object
```

Observations:

1. All the data are in their proper type, no need for any changes.

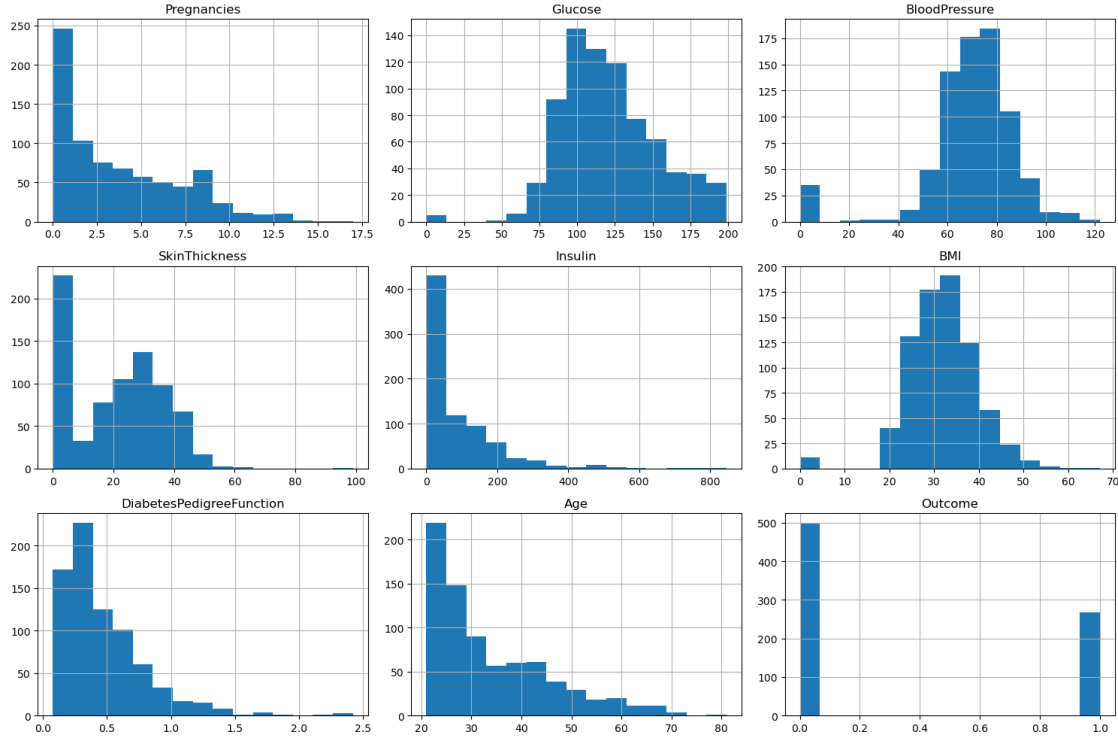
```
[4]: print(df.head())
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

1. Univariate analysis focuses on examining the distribution and characteristics of a single variable. It helps in understanding the central tendency, spread, and shape of the distribution of each variable.
2. Pairwise analysis, on the other hand, explores the relationships between pairs of variables. Pairwise analysis is particularly useful for identifying potential predictor variables or outcomes. It helps in identifying correlations, associations, or patterns between two variables.

```
[5]: # Univariate visualization - Histogram for each feature
df.hist(bins=15, figsize=(15, 10), layout=(3, 3))
plt.tight_layout()
plt.show()
```



Pregnancies: The distribution is heavily right-skewed, with most values concentrated towards 0, indicating that a large portion of the data has few or no pregnancies.

Glucose: The distribution appears to be roughly normal, with the majority of values clustered around the mean (potentially around 120).

BloodPressure: The distribution is slightly right-skewed, with a peak around 70-80 mm Hg, which is within the normal range for diastolic blood pressure.

SkinThickness: The distribution is heavily right-skewed, with a large number of zero or very low values, indicating that many individuals have minimal skin thickness measurements.

Insulin: The distribution is heavily right-skewed, with a large number of zero or very low values, suggesting that many individuals have low or no insulin levels measured.

BMI: The distribution appears to be roughly normal, with the majority of values falling within the overweight or obese range (25-35), potentially indicating a correlation between higher BMI and the presence of diabetes in the dataset.

DiabetesPedigreeFunction: The distribution is heavily right-skewed, with most values concentrated towards 0, indicating that a large portion of the data has low diabetes pedigree function values (lower likelihood of diabetes based on family history).

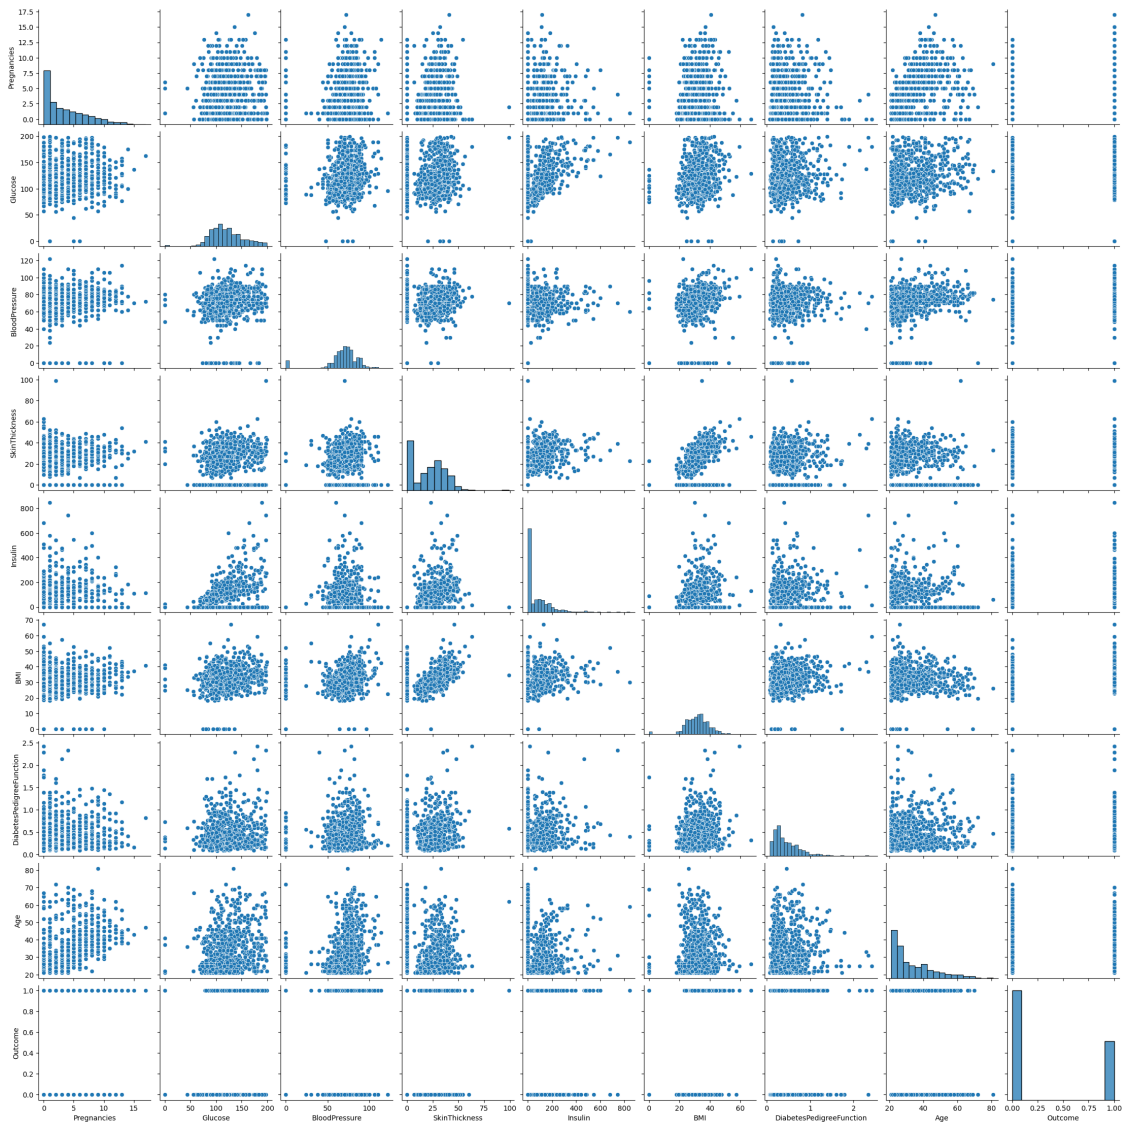
Age: The distribution appears to be roughly normal, with a peak around 30-40 years of age, suggesting that the dataset covers a range of ages.

Outcome: The distribution is bimodal, with two distinct peaks representing the two classes (0 for no diabetes, 1 for diabetes). The smaller peak indicates that there are fewer instances of diabetes.

in the dataset compared to non-diabetes cases.

```
[6]: # Pairwise visualization - Pairplot to observe relationships between features
sns.pairplot(df)
plt.show()
```

```
C:\Users\nav\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\nav\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\nav\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\nav\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\nav\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\nav\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\nav\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\nav\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\nav\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
```

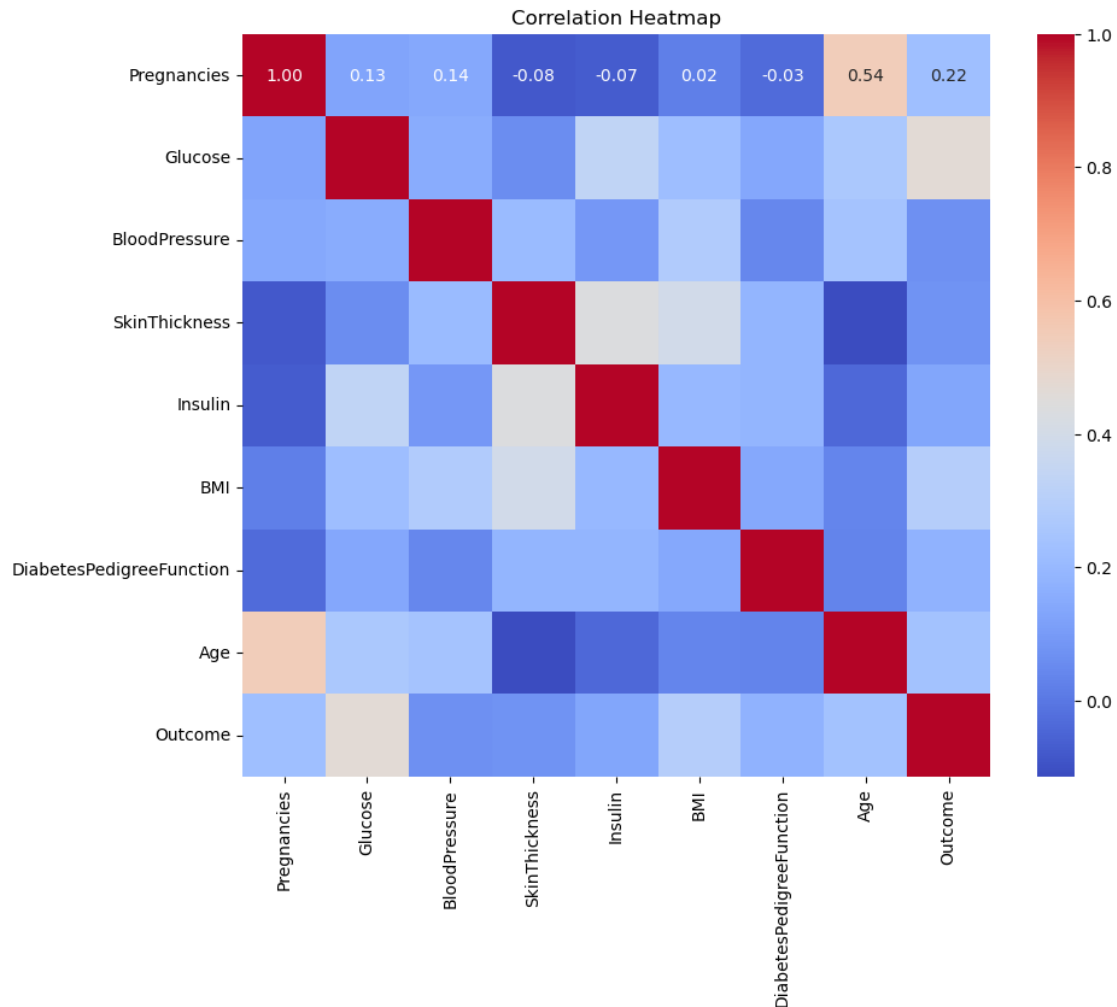


Observations:

1. From the pariplots, we can infer that there is positive correlation between Glucose and Insulin.
2. SkinThickness also seems to increase with an increase in BMI.
3. Blood Pressure does not seem to have any correlation with other features.
4. The other plots do not seem to indicate any correlation.

```
[7]: corr_matrix = df.corr()

# Plot heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap')
plt.show()
```



0.1 Data Cleaning

```
[8]: # Check for initial zero values
zero_counts = (df == 0).sum()
print("Percentage of zero in each column:\n", zero_counts / len(df) * 100)

# Check for initial missing values
initial_missing_values = df.isnull().sum()
print("Initial missing values in each column:\n", initial_missing_values /
      ↪ len(df) * 100)

# Replace zeroes with NaN in specified columns
columns_with_invalid_zeros = ['Glucose', 'BloodPressure', 'SkinThickness',
      ↪ 'Insulin', 'BMI']
```

```
df[columns_with_invalid_zeros] = df[columns_with_invalid_zeros].replace(0, np.
    ↪nan)

# Show the count and percentage of missing values after replacement
missing_values_after_replacement = df.isnull().sum()
missing_values_percentage = (missing_values_after_replacement / len(df)) * 100

print("\nPercentage of missing values in each column after replacing zeroes_
    ↪with NaN:\n", missing_values_percentage)
```

Percentage of zero in each column:

Pregnancies	14.453125
Glucose	0.651042
BloodPressure	4.557292
SkinThickness	29.557292
Insulin	48.697917
BMI	1.432292
DiabetesPedigreeFunction	0.000000
Age	0.000000
Outcome	65.104167

dtype: float64

Initial missing values in each column:

Pregnancies	0.0
Glucose	0.0
BloodPressure	0.0
SkinThickness	0.0
Insulin	0.0
BMI	0.0
DiabetesPedigreeFunction	0.0
Age	0.0
Outcome	0.0

dtype: float64

Percentage of missing values in each column after replacing zeroes with NaN:

Pregnancies	0.000000
Glucose	0.651042
BloodPressure	4.557292
SkinThickness	29.557292
Insulin	48.697917
BMI	1.432292
DiabetesPedigreeFunction	0.000000
Age	0.000000
Outcome	0.000000

dtype: float64

Observation:

1. I will impute 'Glucose', 'BloodPressure', 'BMI' as they have low percentages of missing values.

2. I can either impute or delete the 'SkinThickness' and 'Insulin' columns as they have high an
3. I left out pregnancies and outcome column while replacing '0' with NAN because those columns

```
[9]: # Calculate the median for each column
medians = df.median()

# Fill missing values with the median of each column
df = df.fillna(medians)

# Calculate the percentage of missing values after imputation
missing_values_percentage = (df.isnull().sum() / len(df)) * 100
print("\nPercentage of missing values in each column after imputing:\n",
      ↪missing_values_percentage)
```

Percentage of missing values in each column after imputing:

Pregnancies	0.0
Glucose	0.0
BloodPressure	0.0
SkinThickness	0.0
Insulin	0.0
BMI	0.0
DiabetesPedigreeFunction	0.0
Age	0.0
Outcome	0.0

dtype: float64

0.2 Partition the data and prepare it.

e).

```
[10]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder

# Partition the data into train/test sets
X = df.drop(columns='Outcome')
y = df['Outcome']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=42, stratify=y)

# Normalize numeric data
scaler = StandardScaler()
scaler.fit(X_train)

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

Explanation

1. There are no categorical variables that needs to be encoded here. There are some columns that
2. normalization has been done using standardScaler from sklearn.
$$x = (x - \text{mean}(x)) / \text{standard_deviation}(x)$$
3. fit_transform() computes and transforms the training data, while transform() uses the previ

0.2.1 The documentation of the SVM algorithm to identify the hyperparametined.

```
[11]: from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import classification_report, accuracy_score
import pickle

# Instantiate the SVM model
svc = SVC(random_state=42)

svm = SVC()
svm.fit(X_train, y_train)
y_pred = svm.predict(X_test)
print('Accuracy here: ', svm.score(X_test, y_test))

# Define hyperparameters to tune
param_grid = {
    'C': [0.1, 1, 10, 100], # Tuning best regularization parameter
    'kernel': ['linear', 'rbf', 'sigmoid'] # Kernel to be used. Using kernels
    ↪as they might capture non linear relation if it exists
}

# Use GridSearchCV to tune hyperparameters
grid = GridSearchCV(svc, param_grid, cv=6, scoring='precision', n_jobs=-1,
    ↪return_train_score=True)
grid.fit(X_train, y_train)

# Display the best model/parameters and the best score obtained
print("Best parameters found: ", grid.best_params_)
print("Best score of GridSearch: ", grid.best_score_)

# Testing the performance on test set
best_model = grid.best_estimator_
test_score = best_model.score(X_test, y_test)
print(f"Test set precision of GridSearch: {test_score:.3f}")

# Evaluate the model on the test set
y_pred = grid.predict(X_test)
print("Test Set Accuracy of the Model: ", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```

# Extract results from GridSearchCV
results = grid.cv_results_

# Print mean train and test scores along with standard deviation for each
↳ parameter combination
for mean_train_score, std_train_score, mean_test_score, std_test_score, params_
↳ in zip(
    results['mean_train_score'], results['std_train_score'],
    results['mean_test_score'], results['std_test_score'],
    results['params']):
    print(f"Params: {params}")
    print(f"Mean train score: {mean_train_score:.3f} (std: {std_train_score:.
↳ 3f})")
    print(f"Mean test score: {mean_test_score:.3f} (std: {std_test_score:.
↳ 3f})\n")

# Save the best model and grid search results
with open('best_model.pkl', 'wb') as f:
    pickle.dump(best_model, f)
with open('grid_search_results.pkl', 'wb') as f:
    pickle.dump(grid, f)

```

Accuracy here: 0.7402597402597403

Best parameters found: {'C': 0.1, 'kernel': 'rbf'}

Best score of GridSearch: 0.7504563492063493

Test set precision of GridSearch: 0.734

Test Set Accuracy of the Model: 0.7337662337662337

	precision	recall	f1-score	support
0	0.74	0.90	0.81	100
1	0.70	0.43	0.53	54
accuracy			0.73	154
macro avg	0.72	0.66	0.67	154
weighted avg	0.73	0.73	0.71	154

Params: {'C': 0.1, 'kernel': 'linear'}

Mean train score: 0.758 (std: 0.011)

Mean test score: 0.741 (std: 0.047)

Params: {'C': 0.1, 'kernel': 'rbf'}

Mean train score: 0.828 (std: 0.014)

Mean test score: 0.750 (std: 0.039)

Params: {'C': 0.1, 'kernel': 'sigmoid'}

Mean train score: 0.739 (std: 0.010)
Mean test score: 0.740 (std: 0.049)

Params: {'C': 1, 'kernel': 'linear'}
Mean train score: 0.756 (std: 0.010)
Mean test score: 0.738 (std: 0.036)

Params: {'C': 1, 'kernel': 'rbf'}
Mean train score: 0.841 (std: 0.015)
Mean test score: 0.723 (std: 0.058)

Params: {'C': 1, 'kernel': 'sigmoid'}
Mean train score: 0.573 (std: 0.021)
Mean test score: 0.618 (std: 0.050)

Params: {'C': 10, 'kernel': 'linear'}
Mean train score: 0.756 (std: 0.011)
Mean test score: 0.739 (std: 0.049)

Params: {'C': 10, 'kernel': 'rbf'}
Mean train score: 0.890 (std: 0.009)
Mean test score: 0.660 (std: 0.083)

Params: {'C': 10, 'kernel': 'sigmoid'}
Mean train score: 0.527 (std: 0.017)
Mean test score: 0.549 (std: 0.088)

Params: {'C': 100, 'kernel': 'linear'}
Mean train score: 0.756 (std: 0.011)
Mean test score: 0.739 (std: 0.049)

Params: {'C': 100, 'kernel': 'rbf'}
Mean train score: 0.973 (std: 0.007)
Mean test score: 0.567 (std: 0.029)

Params: {'C': 100, 'kernel': 'sigmoid'}
Mean train score: 0.538 (std: 0.029)
Mean test score: 0.523 (std: 0.090)

```
[12]: import pickle
      from sklearn.metrics import classification_report, accuracy_score

      # Load the best model
      with open('best_model.pkl', 'rb') as f:
          best_model = pickle.load(f)
```

```

# Load the grid search results
with open('grid_search_results.pkl', 'rb') as f:
    grid = pickle.load(f)

# Testing the performance on test set
test_score = best_model.score(X_test, y_test)
print(f"Test set precision of GridSearch: {test_score:.3f}")

# Evaluate the model on the test set
y_pred = grid.predict(X_test)
print("Test Set Accuracy of the Model: ", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))

# Extract results from GridSearchCV
results = grid.cv_results_

# Print mean train and test scores along with standard deviation for each
# parameter combination
for mean_train_score, std_train_score, mean_test_score, std_test_score, params_
    in zip(
        results['mean_train_score'], results['std_train_score'],
        results['mean_test_score'], results['std_test_score'],
        results['params']):
    print(f"Params: {params}")
    print(f"Mean train score: {mean_train_score:.3f} (std: {std_train_score:.
3f})")
    print(f"Mean test score: {mean_test_score:.3f} (std: {std_test_score:.
3f})\n")

```

Test set precision of GridSearch: 0.734

Test Set Accuracy of the Model: 0.7337662337662337

	precision	recall	f1-score	support
0	0.74	0.90	0.81	100
1	0.70	0.43	0.53	54
accuracy			0.73	154
macro avg	0.72	0.66	0.67	154
weighted avg	0.73	0.73	0.71	154

Params: {'C': 0.1, 'kernel': 'linear'}

Mean train score: 0.758 (std: 0.011)

Mean test score: 0.741 (std: 0.047)

Params: {'C': 0.1, 'kernel': 'rbf'}

Mean train score: 0.828 (std: 0.014)

Mean test score: 0.750 (std: 0.039)

Params: {'C': 0.1, 'kernel': 'sigmoid'}
Mean train score: 0.739 (std: 0.010)
Mean test score: 0.740 (std: 0.049)

Params: {'C': 1, 'kernel': 'linear'}
Mean train score: 0.756 (std: 0.010)
Mean test score: 0.738 (std: 0.036)

Params: {'C': 1, 'kernel': 'rbf'}
Mean train score: 0.841 (std: 0.015)
Mean test score: 0.723 (std: 0.058)

Params: {'C': 1, 'kernel': 'sigmoid'}
Mean train score: 0.573 (std: 0.021)
Mean test score: 0.618 (std: 0.050)

Params: {'C': 10, 'kernel': 'linear'}
Mean train score: 0.756 (std: 0.011)
Mean test score: 0.739 (std: 0.049)

Params: {'C': 10, 'kernel': 'rbf'}
Mean train score: 0.890 (std: 0.009)
Mean test score: 0.660 (std: 0.083)

Params: {'C': 10, 'kernel': 'sigmoid'}
Mean train score: 0.527 (std: 0.017)
Mean test score: 0.549 (std: 0.088)

Params: {'C': 100, 'kernel': 'linear'}
Mean train score: 0.756 (std: 0.011)
Mean test score: 0.739 (std: 0.049)

Params: {'C': 100, 'kernel': 'rbf'}
Mean train score: 0.973 (std: 0.007)
Mean test score: 0.567 (std: 0.029)

Params: {'C': 100, 'kernel': 'sigmoid'}
Mean train score: 0.538 (std: 0.029)
Mean test score: 0.523 (std: 0.090)

Reference:

1. <https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html>
2. https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html

Observations:

1. When GridSearchCV is used it automatically uses cross validation, in here cv = 5 means 5 folds

2. Out of all actual class 0 instances, 90% were correctly identified as 'No diabetes'.
3. Out of all actual instances of "Diabetes," only 43% were correctly identified by the model.
4. Out of all instances predicted as "No Diabetes," 74% were actually "No Diabetes."
5. Out of all instances predicted as "Diabetes," 70% were actually "Diabetes."

Bias and Variance

1. The model may be overfitting as the training accuracy is high and the test accuracy has dropped.
2. The model achieved an accuracy of 75% in the training. The 75% suggests that while the predicted values are close to the actual values, there is still a significant difference.
3. The model achieved an accuracy of 73% in the test set meaning that whatever the model has learned from the training set, it is not perfectly generalizable to new, unseen data.

Model

1. Given its performance on training and test set, the classifier seems to be doing a decent job.
2. As we only have an accuracy of 73%, the model can help medical industry as subordinate in assisting diagnosis.

[]: